

Multivariate Analysis for the Behavioral Sciences,
Second Edition (Chapman and Hall/CRC, 2019)

Examples of Chapter 8:
Analysis of Longitudinal Data I: Graphical Displays
and Summary Measure Approach

Kimmo Vehkalahti and Brian S. Everitt

1 October 2018

Contents

Examples	2
Table 8.1: BPRS Measurements from 40 Subjects	2
Figure 8.1	4
Figure 8.2	5
Figure 8.3	6
Figure 8.4	8
Figure 8.5	9
Figure 8.6	11
Table 8.3	12
Table 8.4	13
Table 8.5: Pain Scores from 83 Women in Labor: 43 Subjects in Group 1 (Medication) and 40 Subjects in Group 2 (Placebo)	14
Table 8.6	15

Examples

Table 8.1: BPRS Measurements from 40 Subjects

```
library(tidyr); library(dplyr); library(ggplot2)

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
BPRS <- read.table("data/BPRS.txt", sep = ' ', header = TRUE)
BPRS <- within(BPRS, {
  treatment <- factor(treatment)
  subject <- factor(subject)
})

glimpse(BPRS)

## Rows: 40
## Columns: 11
## $ treatment <fct> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ~
## $ subject <fct> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 1~
## $ week0 <int> 42, 58, 54, 55, 72, 48, 71, 30, 41, 57, 30, 55, 36, 38, 66, ~
## $ week1 <int> 36, 68, 55, 77, 75, 43, 61, 36, 43, 51, 34, 52, 32, 35, 68, ~
## $ week2 <int> 36, 61, 41, 49, 72, 41, 47, 38, 39, 51, 34, 49, 36, 36, 65, ~
## $ week3 <int> 43, 55, 38, 54, 65, 38, 30, 38, 35, 55, 41, 54, 31, 34, 49, ~
## $ week4 <int> 41, 43, 43, 56, 50, 36, 27, 31, 28, 53, 36, 48, 25, 25, 36, ~
## $ week5 <int> 40, 34, 28, 50, 39, 29, 40, 26, 22, 43, 36, 43, 25, 27, 32, ~
## $ week6 <int> 38, 28, 29, 47, 32, 33, 30, 26, 20, 43, 38, 37, 21, 25, 27, ~
## $ week7 <int> 47, 28, 25, 42, 38, 27, 31, 25, 23, 39, 36, 36, 19, 26, 30, ~
## $ week8 <int> 51, 28, 24, 46, 32, 25, 31, 24, 21, 32, 36, 31, 22, 26, 37, ~

head(BPRS)

##   treatment subject week0 week1 week2 week3 week4 week5 week6 week7 week8
## 1          1         1    42    36    36    43    41    40    38    47    51
## 2          1         2    58    68    61    55    43    34    28    28    28
## 3          1         3    54    55    41    38    43    28    29    25    24
## 4          1         4    55    77    49    54    56    50    47    42    46
## 5          1         5    72    75    72    65    50    39    32    38    32
## 6          1         6    48    43    41    38    36    29    33    27    25

tail(BPRS)

##   treatment subject week0 week1 week2 week3 week4 week5 week6 week7 week8
## 35          2        15    40    36    55    55    42    30    26    30    37
## 36          2        16    54    45    35    27    25    22    22    22    22
## 37          2        17    33    41    30    32    46    43    43    43    43
## 38          2        18    28    30    29    33    30    26    36    33    30
```

```
## 39      2      19      52      43      26      27      24      32      21      21      21
## 40      2      20      47      36      32      29      25      23      23      23      23
```

Convert data to long form:

```
BPRSL <- gather(BPRS, key = weeks, value = bprs, week0:week8) %>%
  mutate(week = as.integer(substr(weeks, 5, 5)))
```

```
glimpse(BPRSL)
```

```
## Rows: 360
```

```
## Columns: 5
```

```
## $ treatment <fct> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ~
```

```
## $ subject   <fct> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 1~
```

```
## $ weeks     <chr> "week0", "week0", "week0", "week0", "week0", "week0", "week0~
```

```
## $ bprs      <int> 42, 58, 54, 55, 72, 48, 71, 30, 41, 57, 30, 55, 36, 38, 66, ~
```

```
## $ week      <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ~
```

```
head(BPRSL)
```

```
##   treatment subject weeks bprs week
```

```
## 1          1       1 week0  42    0
```

```
## 2          1       2 week0  58    0
```

```
## 3          1       3 week0  54    0
```

```
## 4          1       4 week0  55    0
```

```
## 5          1       5 week0  72    0
```

```
## 6          1       6 week0  48    0
```

```
tail(BPRSL)
```

```
##   treatment subject weeks bprs week
```

```
## 355        2      15 week8  37    8
```

```
## 356        2      16 week8  22    8
```

```
## 357        2      17 week8  43    8
```

```
## 358        2      18 week8  30    8
```

```
## 359        2      19 week8  21    8
```

```
## 360        2      20 week8  23    8
```

Figure 8.1

```
p1 <- ggplot(BPRSL, aes(x = week, y = bprs, linetype = subject))
p2 <- p1 + geom_line() + scale_linetype_manual(values = rep(1:10, times=4))
p3 <- p2 + facet_grid(. ~ treatment, labeller = label_both)
p4 <- p3 + theme_bw() + theme(legend.position = "none")
p5 <- p4 + theme(panel.grid.minor.y = element_blank())
p6 <- p5 + scale_y_continuous(limits = c(min(BPRSL$bprs), max(BPRSL$bprs)))
p6
```




Figure 8.2

```
# Standardise the scores:
BPRSL <- BPRSL %>%
  group_by(week) %>%
  mutate( stdbprs = (bprs - mean(bprs))/sd(bprs) ) %>%
  ungroup()
glimpse(BPRSL)

## Rows: 360
## Columns: 6
## $ treatment <fct> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ~
## $ subject <fct> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 1~
## $ weeks <chr> "week0", "week0", "week0", "week0", "week0", "week0", "week0~
## $ bprs <int> 42, 58, 54, 55, 72, 48, 71, 30, 41, 57, 30, 55, 36, 38, 66, ~
## $ week <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ~
## $ stdbprs <dbl> -0.4245908, 0.7076513, 0.4245908, 0.4953559, 1.6983632, 0.00~

p1 <- ggplot(BPRSL, aes(x = week, y = stdbprs, linetype = subject))
p2 <- p1 + geom_line() + scale_linetype_manual(values = rep(1:10, times=4))
p3 <- p2 + facet_grid(. ~ treatment, labeller = label_both)
p4 <- p3 + theme_bw() + theme(legend.position = "none")
p5 <- p4 + theme(panel.grid.minor.y = element_blank())
p6 <- p5 + scale_y_continuous(name = "standardized bprs")
p6
```




Figure 8.3

```
# Number of subjects: (would be better to have unique IDs...)
n <- 40
# Make a summary data:
BPRSS <- BPRSL %>%
  group_by(treatment, week) %>%
  summarise( mean=mean(bprs), se=sd(bprs)/sqrt(n) ) %>%
  ungroup()

## `summarise()` has grouped output by 'treatment'. You can override using the `.groups` argument.
glimpse(BPRSS)

## Rows: 18
## Columns: 4
## $ treatment <fct> 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2
## $ week      <int> 0, 1, 2, 3, 4, 5, 6, 7, 8, 0, 1, 2, 3, 4, 5, 6, 7, 8
## $ mean      <dbl> 47.00, 46.80, 43.55, 40.90, 36.60, 32.70, 29.70, 29.80, 29.3~
## $ se        <dbl> 2.150887, 2.454105, 1.899082, 1.776232, 1.546133, 1.231281, ~

p1 <- ggplot(BPRSS, aes(x = week, y = mean, linetype = treatment, shape = treatment))
p2 <- p1 + geom_line() + scale_linetype_manual(values = c(1,2))
p3 <- p2 + geom_point(size=3) + scale_shape_manual(values = c(1,2))
p4 <- p3 + geom_errorbar(aes(ymin=mean-se, ymax=mean+se, linetype="1"), width=0.3)
p5 <- p4 + theme_bw() + theme(panel.grid.major = element_blank(), panel.grid.minor = element_blank())
p6 <- p5 + theme(legend.position = c(0.8,0.8))
p7 <- p6 + scale_y_continuous(name = "mean(bprs) +/- se(bprs)")
p7
```




Figure 8.4

```
p1 <- ggplot(BPRSL, aes(x = factor(week), y = bprs, fill = treatment))
p2 <- p1 + geom_boxplot(position = position_dodge(width = 0.9))
p3 <- p2 + theme_bw() + theme(panel.grid.major = element_blank(), panel.grid.minor = element_blank())
p4 <- p3 + theme(legend.position = c(0.8,0.8))
p5 <- p4 + scale_x_discrete(name = "week")
# Black & White version:
#p6 <- p5 + scale_fill_grey(start = 0.5, end = 1)
p5
```




Figure 8.5

```
# Make a summary data of the post treatment weeks (1-8)
BPRSL8S <- BPRSL %>%
  filter(week > 0) %>%
  group_by(treatment, subject) %>%
  summarise( mean=mean(bprs) ) %>%
  ungroup()

## `summarise()` has grouped output by 'treatment'. You can override using the `.groups` argument.
glimpse(BPRSL8S)

## Rows: 40
## Columns: 3
## $ treatment <fct> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ~
## $ subject <fct> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 1~
## $ mean <dbl> 41.500, 43.125, 35.375, 52.625, 50.375, 34.000, 37.125, 30.5~

p1 <- ggplot(BPRSL8S, aes(x = treatment, y = mean))
p2 <- p1 + geom_boxplot()
p3 <- p2 + theme_bw() + theme(panel.grid.major = element_blank(), panel.grid.minor = element_blank())
p4 <- p3 + stat_summary(fun.y = "mean", geom = "point", shape=23, size=4, fill = "white")

## Warning: `fun.y` is deprecated. Use `fun` instead.
p5 <- p4 + scale_y_continuous(name = "mean(bprs), weeks 1-8")
p5
```




Figure 8.6

```
# Remove the outlier:
BPRSL8S1 <- BPRSL8S %>%
  filter(mean < 60)

glimpse(BPRSL8S1)

## Rows: 39
## Columns: 3
## $ treatment <fct> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ~
## $ subject    <fct> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 1~
## $ mean       <dbl> 41.500, 43.125, 35.375, 52.625, 50.375, 34.000, 37.125, 30.5~

p1 <- ggplot(BPRSL8S1, aes(x = treatment, y = mean))
p2 <- p1 + geom_boxplot()
p3 <- p2 + theme_bw() + theme(panel.grid.major = element_blank(), panel.grid.minor = element_blank())
p4 <- p3 + stat_summary(fun.y = "mean", geom = "point", shape=23, size=4, fill = "white")

## Warning: `fun.y` is deprecated. Use `fun` instead.
p5 <- p4 + scale_y_continuous(name = "mean(bprs), weeks 1-8")
p5
```




Table 8.3

```
# Without the outlier, apply Student's t-test, two-sided:

t.test(mean ~ treatment, data = BPRSL8S1, var.equal = TRUE)

##
## Two Sample t-test
##
## data: mean by treatment
## t = 0.52095, df = 37, p-value = 0.6055
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -4.232480 7.162085
## sample estimates:
## mean in group 1 mean in group 2
## 36.16875 34.70395
```

Table 8.4

```
# Add the baseline from the original data as a new variable to the summary data:
```

```
baseline <- BPRS$week0
BPRSL8S2 <- BPRSL8S %>%
  mutate(baseline)
```

```
# Fit the ANCOVA model and see the results:
```

```
fit <- lm(mean ~ baseline + treatment, data = BPRSL8S2)
```

```
summary(fit)
```

```
##
```

```
## Call:
```

```
## lm(formula = mean ~ baseline + treatment, data = BPRSL8S2)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -16.1729  -4.5994   0.1088   4.6703  21.0656
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 13.07897    4.55710   2.870  0.00675 **
## baseline     0.49127    0.08943   5.493 3.05e-06 ***
## treatment2  -0.58879    2.49584  -0.236  0.81480
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## Residual standard error: 7.872 on 37 degrees of freedom
```

```
## Multiple R-squared:  0.4494, Adjusted R-squared:  0.4196
```

```
## F-statistic: 15.1 on 2 and 37 DF, p-value: 1.605e-05
```

```
anova(fit)
```

```
## Analysis of Variance Table
```

```
##
```

```
## Response: mean
```

```
##      Df Sum Sq Mean Sq F value    Pr(>F)
## baseline  1 1868.07 1868.07 30.1437 3.077e-06 ***
## treatment  1   3.45   3.45  0.0557  0.8148
## Residuals 37 2292.97   61.97
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Table 8.5: Pain Scores from 83 Women in Labor: 43 Subjects in Group 1 (Medication) and 40 Subjects in Group 2 (Placebo)

```
PAIN <- read.table("data/pain.txt", header = TRUE, sep = "\t")

PAIN <- within(PAIN, {
  group <- factor(group)
  id <- factor(id)
})

glimpse(PAIN)

## Rows: 83
## Columns: 9
## $ group <fct> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
## $ id <fct> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 1~
## $ m0 <dbl> 0.0, 0.0, 38.0, 6.0, 19.0, 7.0, 44.0, 1.0, 24.5, 1.0, 35.5, 0.0, ~
## $ m30 <dbl> 0.0, 0.0, 5.0, 48.0, 5.0, 0.0, 42.0, 0.0, 35.0, 30.5, 44.5, 0.0, ~
## $ m60 <dbl> 0.0, 0.0, 1.0, 85.0, NA, 0.0, 42.0, 0.0, 13.0, 81.5, 55.0, 0.0, ~
## $ m90 <dbl> 0.0, 0.0, 1.0, 0.0, NA, 0.0, 45.0, 0.0, NA, 67.5, 69.0, 0.0, 24.~
## $ m120 <dbl> NA, 2.5, 0.0, 0.0, NA, NA, NA, 0.0, NA, 98.5, 72.5, 0.0, 29.0, 1~
## $ m150 <dbl> NA, 2.3, 5.0, NA, NA, NA, NA, 6.0, NA, 97.0, 39.5, 0.0, 45.0, NA~
## $ m180 <dbl> NA, 14.0, NA, NA, NA, NA, NA, 24.0, NA, NA, 26.0, 0.0, 91.0, NA,~
```

Table 8.6

```
# Convert data to long form:
PAINL <- gather(PAIN, key = mins, value = pain, -group, -id) %>%
  mutate(mins = as.integer(substr(mins,2,4)))

glimpse(PAINL)

## Rows: 581
## Columns: 4
## $ group <fct> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
## $ id <fct> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 1~
## $ mins <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ pain <dbl> 0.0, 0.0, 38.0, 6.0, 19.0, 7.0, 44.0, 1.0, 24.5, 1.0, 35.5, 0.0,~

# (1) Make a summary data of group means,
# removing subjects with any missing values:

PAINLS1 <- PAINL %>%
  group_by(group, id) %>%
  summarise( mean=mean(pain) ) %>%
  ungroup()

## `summarise()` has grouped output by 'group'. You can override using the `.groups` argument.
glimpse(PAINLS1)

## Rows: 83
## Columns: 3
## $ group <fct> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
## $ id <fct> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 1~
## $ mean <dbl> NA, 2.6857143, NA, NA, NA, NA, NA, NA, 4.4285714, NA, NA, 48.8571429~

# (2) Make a summary data of group means,
# now using the mean of available responses for each subject:

PAINLS2 <- PAINL %>%
  group_by(group, id) %>%
  summarise( mean=mean(pain, na.rm = TRUE) ) %>%
  ungroup()

## `summarise()` has grouped output by 'group'. You can override using the `.groups` argument.
glimpse(PAINLS2)

## Rows: 83
## Columns: 3
## $ group <fct> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
## $ id <fct> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 1~
## $ mean <dbl> 0.0000000, 2.6857143, 8.3333333, 27.8000000, 12.0000000, 1.75000~

# Compare the t-test results:

t.test(mean ~ group, data = PAINLS1, var.equal = TRUE)

##
## Two Sample t-test
```

```
##
## data: mean by group
## t = -3.8465, df = 32, p-value = 0.000538
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -45.97397 -14.14032
## sample estimates:
## mean in group 1 mean in group 2
## 16.74286 46.80000
t.test(mean ~ group, data = PAINLS2, var.equal = TRUE)

##
## Two Sample t-test
##
## data: mean by group
## t = -4.1455, df = 81, p-value = 8.293e-05
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -33.58920 -11.80261
## sample estimates:
## mean in group 1 mean in group 2
## 18.34181 41.03771
```