

*Multivariate Analysis for the Behavioral Sciences,*  
Second Edition (Chapman and Hall/CRC, 2019)

**Examples of Chapter 10:**  
**Analysis of Longitudinal Data III:**  
**Non-Normal Responses**

*Kimmo Vehkalahti and Brian S. Everitt*

*22 November 2018*

**Contents**

<b>Examples</b>	<b>2</b>
<b>Beat the Blues Revisited</b>	<b>2</b>
Table 10.3 . . . . .	4
Table 10.4 . . . . .	5
<b>Table 10.1: Respiratory Disorder Data</b>	<b>6</b>
Figure 10.2 . . . . .	8
Table 10.5 . . . . .	10
Table 10.6 . . . . .	11
Table 10.7 . . . . .	12
Table 10.10 . . . . .	14
<b>Table 10.2: Data from a Clinical Trial of Patients Suffering from Epilepsy</b>	<b>16</b>
Figure 10.3 . . . . .	18
Table 10.8 . . . . .	20
Table 10.9 . . . . .	21
Table 10.11 . . . . .	23
Table 10.12 . . . . .	24

## Examples

### Beat the Blues Revisited

Using the Beat the Blues data introduced in **Chapter 9**:

```
library("tidyr"); library("dplyr")

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

BtB <- read.table("data/BtB.txt", header = TRUE, sep = '\t')

# to make sure that the factor levels are logical (esp. Treatment):
BtB <- within(BtB, {
  Drug <- factor(Drug, levels=c("No", "Yes")) # default
  Length <- factor(Length, levels=c("<6m", ">6m")) # default
  Treatment <- factor(Treatment, levels=c("TAU", "BtheB")) # NOT default!
})

glimpse(BtB); head(BtB); tail(BtB)
```

```
## Observations: 100
## Variables: 9
## $ Subject <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 1...
## $ Drug <fct> No, Yes, Yes, No, Yes, Yes, Yes, No, Yes, Yes, No, Y...
## $ Length <fct> >6m, >6m, <6m, >6m, >6m, <6m, <6m, >6m, <6m, >6m, >6...
## $ Treatment <fct> TAU, BtheB, TAU, BtheB, BtheB, BtheB, TAU, TAU, Bthe...
## $ BDIpre <int> 29, 32, 25, 21, 26, 7, 17, 20, 18, 20, 30, 49, 26, 3...
## $ BDI2m <int> 2, 16, 20, 17, 23, 0, 7, 20, 13, 5, 32, 35, 27, 26, ...
## $ BDI4m <int> 2, 24, NA, 16, NA, 0, 7, 21, 14, 5, 24, NA, 23, 36, ...
## $ BDI6m <int> NA, 17, NA, 10, NA, 0, 3, 19, 20, 8, 12, NA, NA, 27,...
## $ BDI8m <int> NA, 20, NA, 9, NA, 0, 7, 13, 11, 12, 2, NA, NA, 22, ...
```

```
##   Subject Drug Length Treatment BDIpre BDI2m BDI4m BDI6m BDI8m
## 1      1   No   >6m      TAU      29      2      2    NA    NA
## 2      2  Yes   >6m    BtheB      32     16     24     17     20
## 3      3  Yes   <6m      TAU      25     20     NA     NA     NA
## 4      4   No   >6m    BtheB      21     17     16     10      9
## 5      5  Yes   >6m    BtheB      26     23     NA     NA     NA
## 6      6  Yes   <6m    BtheB       7      0      0      0      0

##   Subject Drug Length Treatment BDIpre BDI2m BDI4m BDI6m BDI8m
## 95      95   No   >6m    BtheB      16     11      4      2      3
## 96      96  Yes   >6m    BtheB      16     16     10     10      8
## 97      97  Yes   <6m      TAU      28     NA     NA     NA     NA
## 98      98   No   >6m    BtheB      11     22      9     11     11
## 99      99   No   <6m      TAU      13      5      5      0      6
```

```
## 100      100 Yes    <6m      TAU      43      NA      NA      NA      NA
```

```
# Convert data to long form for the analyses, adding Time:
```

```
BtBL <- gather(BtB, key = Visit, value = BDI, BDI2m, BDI4m, BDI6m, BDI8m) %>%
  mutate(Time = as.integer(substr(Visit, 4, 4)))
```

```
glimpse(BtBL); head(BtBL); tail(BtBL)
```

```
## Observations: 400
```

```
## Variables: 8
```

```
## $ Subject <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 1...
```

```
## $ Drug <fct> No, Yes, Yes, No, Yes, Yes, Yes, No, Yes, Yes, No, Y...
```

```
## $ Length <fct> >6m, >6m, <6m, >6m, >6m, <6m, <6m, >6m, <6m, >6m, >6...
```

```
## $ Treatment <fct> TAU, BtheB, TAU, BtheB, BtheB, BtheB, TAU, TAU, Bthe...
```

```
## $ BDIpre <int> 29, 32, 25, 21, 26, 7, 17, 20, 18, 20, 30, 49, 26, 3...
```

```
## $ Visit <chr> "BDI2m", "BDI2m", "BDI2m", "BDI2m", "BDI2m", "BDI2m"...
```

```
## $ BDI <int> 2, 16, 20, 17, 23, 0, 7, 20, 13, 5, 32, 35, 27, 26, ...
```

```
## $ Time <int> 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2...
```

```
## Subject Drug Length Treatment BDIpre Visit BDI Time
```

```
## 1      1 No    >6m      TAU      29 BDI2m  2    2
```

```
## 2      2 Yes   >6m      BtheB     32 BDI2m 16    2
```

```
## 3      3 Yes   <6m      TAU      25 BDI2m 20    2
```

```
## 4      4 No    >6m      BtheB     21 BDI2m 17    2
```

```
## 5      5 Yes   >6m      BtheB     26 BDI2m 23    2
```

```
## 6      6 Yes   <6m      BtheB      7 BDI2m  0    2
```

```
## Subject Drug Length Treatment BDIpre Visit BDI Time
```

```
## 395     95 No    >6m      BtheB     16 BDI8m  3    8
```

```
## 396     96 Yes   >6m      BtheB     16 BDI8m  8    8
```

```
## 397     97 Yes   <6m      TAU      28 BDI8m NA    8
```

```
## 398     98 No    >6m      BtheB     11 BDI8m 11    8
```

```
## 399     99 No    <6m      TAU      13 BDI8m  6    8
```

```
## 400    100 Yes   <6m      TAU      43 BDI8m NA    8
```

```
# important to order the data with respect to the subjects:
```

```
BtBLSORT <- BtBL %>%
```

```
  arrange(Subject)
```

```
head(BtBLSORT); tail(BtBLSORT)
```

```
## Subject Drug Length Treatment BDIpre Visit BDI Time
```

```
## 1      1 No    >6m      TAU      29 BDI2m  2    2
```

```
## 2      1 No    >6m      TAU      29 BDI4m  2    4
```

```
## 3      1 No    >6m      TAU      29 BDI6m NA    6
```

```
## 4      1 No    >6m      TAU      29 BDI8m NA    8
```

```
## 5      2 Yes   >6m      BtheB     32 BDI2m 16    2
```

```
## 6      2 Yes   >6m      BtheB     32 BDI4m 24    4
```

```
## Subject Drug Length Treatment BDIpre Visit BDI Time
```

```
## 395     99 No    <6m      TAU      13 BDI6m  0    6
```

```
## 396     99 No    <6m      TAU      13 BDI8m  6    8
```

```
## 397    100 Yes   <6m      TAU      43 BDI2m NA    2
```

```
## 398    100 Yes   <6m      TAU      43 BDI4m NA    4
```

```
## 399    100 Yes   <6m      TAU      43 BDI6m NA    6
```

```
## 400    100 Yes   <6m      TAU      43 BDI8m NA    8
```

Table 10.3

```
# install.packages("gee")
library("gee")

BtB_gee1 <- gee(BDI ~ BDIPre + Time + Treatment + Drug + Length, id = Subject,
               data = BtBLSORT, family = gaussian, corstr = "independence")

## Beginning Cgee S-function, @(#) geeformula.q 4.13 98/01/27
## running glm to get initial regression estimate

##      (Intercept)          BDIPre          Time TreatmentBtheB          DrugYes
##      7.8830747          0.5723729         -0.9608077         -3.3539662         -3.5460058
##      Length>6m
##      1.7530800

summary(BtB_gee1)

##
## GEE: GENERALIZED LINEAR MODELS FOR DEPENDENT DATA
## gee S-function, version 4.13 modified 98/01/27 (1998)
##
## Model:
## Link: Identity
## Variance to Mean Relation: Gaussian
## Correlation Structure: Independent
##
## Call:
## gee(formula = BDI ~ BDIPre + Time + Treatment + Drug + Length,
##      id = Subject, data = BtBLSORT, family = gaussian, corstr = "independence")
##
## Summary of Residuals:
##      Min          1Q          Median          3Q          Max
## -24.20158432  -5.31202378   0.01101526   5.29503741  27.77789553
##
##
## Coefficients:
##      Estimate Naive S.E.   Naive z Robust S.E.  Robust z
## (Intercept)   7.8830747 1.78048908  4.427477  2.19973944  3.583640
## BDIPre        0.5723729 0.05486079 10.433188  0.08853253  6.465114
## Time         -0.9608077 0.23263437 -4.130119  0.17688635 -5.431780
## TreatmentBtheB -3.3539662 1.09831939 -3.053726  1.71390982 -1.956909
## DrugYes       -3.5460058 1.14469086 -3.097785  1.73069664 -2.048889
## Length>6m     1.7530800 1.10849861  1.581490  1.41954159  1.234962
##
## Estimated Scale Parameter: 74.8854
## Number of Iterations: 1
##
## Working Correlation
##      [,1] [,2] [,3] [,4]
## [1,]    1    0    0    0
## [2,]    0    1    0    0
## [3,]    0    0    1    0
## [4,]    0    0    0    1
```

Table 10.4

```
BtB_gee2 <- gee(BDI ~ BDIPre + Time + Treatment + Drug + Length, id = Subject,
               data = BtBLSORT, family = gaussian, corstr = "exchangeable")

## Beginning Cgee S-function, @(#) geeformula.q 4.13 98/01/27
## running glm to get initial regression estimate

##      (Intercept)          BDIPre          Time TreatmentBtheB          DrugYes
##      7.8830747          0.5723729         -0.9608077         -3.3539662         -3.5460058
##      Length>6m
##      1.7530800

summary(BtB_gee2)

##
## GEE: GENERALIZED LINEAR MODELS FOR DEPENDENT DATA
## gee S-function, version 4.13 modified 98/01/27 (1998)
##
## Model:
## Link:                               Identity
## Variance to Mean Relation: Gaussian
## Correlation Structure:               Exchangeable
##
## Call:
## gee(formula = BDI ~ BDIPre + Time + Treatment + Drug + Length,
##      id = Subject, data = BtBLSORT, family = gaussian, corstr = "exchangeable")
##
## Summary of Residuals:
##      Min      1Q      Median      3Q      Max
## -25.4478843  -6.3276726  -0.8152833   4.3622258  25.4078115
##
##
## Coefficients:
##              Estimate Naive S.E.    Naive z Robust S.E.    Robust z
## (Intercept)   5.8855129 2.32380381   2.5327065   2.10712166   2.7931529
## BDIPre         0.6399964 0.08033495   7.9665999   0.07931263   8.0692874
## Time          -0.7078407 0.14254124  -4.9658660   0.15394156  -4.5981134
## TreatmentBtheB -2.3360241 1.72621751  -1.3532617   1.66217026  -1.4054060
## DrugYes        -2.7742506 1.78397557  -1.5550945   1.64824318  -1.6831561
## Length>6m       0.2084783 1.69179766   0.1232288   1.48052530   0.1408137
##
## Estimated Scale Parameter: 77.14393
## Number of Iterations: 5
##
## Working Correlation
##      [,1] [,2] [,3] [,4]
## [1,] 1.0000000 0.6915241 0.6915241 0.6915241
## [2,] 0.6915241 1.0000000 0.6915241 0.6915241
## [3,] 0.6915241 0.6915241 1.0000000 0.6915241
## [4,] 0.6915241 0.6915241 0.6915241 1.0000000
```

## Table 10.1: Respiratory Disorder Data

```
library("tidyr"); library("dplyr")
# install.packages("HSAUR3")
library("HSAUR3")

## Loading required package: tools

data("respiratory", package = "HSAUR3")
RESPL <- respiratory # (data are readily given in long form)
glimpse(RESPL); head(RESPL, n = 5)

## Observations: 555
## Variables: 7
## $ centre    <fct> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
## $ treatment <fct> placebo, placebo, placebo, placebo, placebo, placebo, placebo...
## $ gender    <fct> female, female, female, female, female, female, fema...
## $ age       <dbl> 46, 46, 46, 46, 46, 28, 28, 28, 28, 28, 23, 23, 23, ...
## $ status    <fct> poor, poor, poor, poor, poor, poor, poor, poor, poor, poor...
## $ month     <ord> 0, 1, 2, 3, 4, 0, 1, 2, 3, 4, 0, 1, 2, 3, 4, 0, 1, 2...
## $ subject   <fct> 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 3, 3, 3, 3, 3, 4, 4, 4...

##   centre treatment gender age status month subject
## 1      1 placebo female 46   poor    0      1
## 112    1 placebo female 46   poor    1      1
## 223    1 placebo female 46   poor    2      1
## 334    1 placebo female 46   poor    3      1
## 445    1 placebo female 46   poor    4      1

RESPL <- within(RESPL, {
  levels(gender) <- c("male", "female") # see the original article by Davis (1991)
  status <- as.numeric(status) - 1 # for converting to wide form below
})
head(RESPL, n = 5)

##   centre treatment gender age status month subject
## 1      1 placebo   male 46     0     0      1
## 112    1 placebo   male 46     0     1      1
## 223    1 placebo   male 46     0     2      1
## 334    1 placebo   male 46     0     3      1
## 445    1 placebo   male 46     0     4      1

# convert data to WIDE form:
# (sep: set valid names month0:month4 here; otherwise just 0:4)
RESP <- RESPL %>%
  spread(key = month, value = status, sep = "") %>%
  arrange(subject)

glimpse(RESP); head(RESP, n = 5)

## Observations: 111
## Variables: 10
## $ centre    <fct> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
## $ treatment <fct> placebo, placebo, treatment, placebo, placebo, treat...
## $ gender    <fct> male, male, male, male, female, male, male, male, ma...
## $ age       <dbl> 46, 28, 23, 44, 13, 34, 43, 28, 31, 37, 30, 14, 23, ...
```

```
## $ subject <fct> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 1...
## $ month0 <dbl> 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0, 0...
## $ month1 <dbl> 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0...
## $ month2 <dbl> 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1...
## $ month3 <dbl> 0, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1...
## $ month4 <dbl> 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1...
```

```
##   centre treatment gender age subject month0 month1 month2 month3 month4
## 1     1   placebo   male  46         1      0      0      0      0      0
## 2     1   placebo   male  28         2      0      0      0      0      0
## 3     1 treatment   male  23         3      1      1      1      1      1
## 4     1   placebo   male  44         4      1      1      1      1      0
## 5     1   placebo female  13         5      1      1      1      1      1
```

*# back to long form; leave month0 for baseline:*

```
RESPL <- gather(RESPL, key = months, value = status, month1:month4) %>%
  mutate(month = as.integer(substr(months, 6, 6))) %>%
  rename(baseline = month0) %>%
  arrange(subject)
```

```
glimpse(RESPL); head(RESPL, n = 5)
```

```
## Observations: 444
```

```
## Variables: 9
```

```
## $ centre <fct> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
## $ treatment <fct> placebo, placebo, placebo, placebo, placebo, placebo...
## $ gender <fct> male, male, male, male, male, male, male, male, male...
## $ age <dbl> 46, 46, 46, 46, 28, 28, 28, 28, 23, 23, 23, 23, 44, ...
## $ subject <fct> 1, 1, 1, 1, 2, 2, 2, 2, 3, 3, 3, 3, 4, 4, 4, 4, 5, 5...
## $ baseline <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
## $ months <chr> "month1", "month2", "month3", "month4", "month1", "m...
## $ status <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1...
## $ month <int> 1, 2, 3, 4, 1, 2, 3, 4, 1, 2, 3, 4, 1, 2, 3, 4, 1, 2...
```

```
##   centre treatment gender age subject baseline months status month
## 1     1   placebo   male  46         1      0 month1      0      1
## 2     1   placebo   male  46         1      0 month2      0      2
## 3     1   placebo   male  46         1      0 month3      0      3
## 4     1   placebo   male  46         1      0 month4      0      4
## 5     1   placebo   male  28         2      0 month1      0      1
```

Figure 10.2

```
n = dim(RESP)[1]
RP <- RESP %>%
  group_by(treatment) %>%
  summarise(m0 = 100*sum(month0)/n,
            m1 = 100*sum(month1)/n,
            m2 = 100*sum(month2)/n,
            m3 = 100*sum(month3)/n,
            m4 = 100*sum(month4)/n ) %>%
  select(-treatment) %>%
  ungroup()

RP2 <- cbind(RP[1, ], RP[2, ])
month <- rep(1:5, times = 2)
treatment <- rep(0:1, each = 5)
RPos <- as.data.frame(cbind(treatment, month, t(RP2)))
row.names(RPos) <- 1:10
names(RPos) <- c("treatment", "month", "positive")
RPos$treatment <- factor(RPos$treatment, labels = c("placebo", "treatment"))

library("ggplot2")

p1 <- ggplot(RPos, aes(x = month, y = positive, fill = treatment))
p2 <- p1 + geom_bar(position = "dodge", stat = "identity")
p3 <- p2 + scale_x_continuous(name = "", breaks = 1:5,
                             labels = c("Baseline", "Month 1", "Month 2", "Month 3", "Month 4"))
p4 <- p3 + theme_bw() + theme(panel.grid.major = element_blank(),
                             panel.grid.minor = element_blank())
p5 <- p4 + scale_y_continuous(name = "Positive responses (%)")
p6 <- p5 + scale_fill_grey(start = 0.5, end = 0)
p6
```






Table 10.5

```

resp_gee1 <- gee(status ~ age, data = RESPL, family = binomial, id = subject,
                 corstr = "independence", scale.fix = TRUE, scale.value = 1)

## Beginning Cgee S-function, @(#) geeformula.q 4.13 98/01/27
## running glm to get initial regression estimate
## (Intercept)          age
## 0.64576265 -0.01201525

summary(resp_gee1)

##
## GEE: GENERALIZED LINEAR MODELS FOR DEPENDENT DATA
## gee S-function, version 4.13 modified 98/01/27 (1998)
##
## Model:
## Link:                               Logit
## Variance to Mean Relation: Binomial
## Correlation Structure:      Independent
##
## Call:
## gee(formula = status ~ age, id = subject, data = RESPL, family = binomial,
##      corstr = "independence", scale.fix = TRUE, scale.value = 1)
##
## Summary of Residuals:
##      Min       1Q   Median       3Q      Max
## -0.6256488 -0.5508736  0.3828310  0.4320996  0.5427142
##
##
## Coefficients:
##              Estimate Naive S.E.   Naive z Robust S.E.  Robust z
## (Intercept) 0.64576265 0.255300707  2.529420  0.39567302  1.632061
## age        -0.01201525 0.007061029 -1.701628  0.01173464 -1.023912
##
## Estimated Scale Parameter: 1
## Number of Iterations: 1
##
## Working Correlation
##      [,1] [,2] [,3] [,4]
## [1,]    1    0    0    0
## [2,]    0    1    0    0
## [3,]    0    0    1    0
## [4,]    0    0    0    1

```

Table 10.6

```

resp_gee2 <- gee(status ~ age, data = RESPL, family = binomial, id = subject,
                 corstr = "exchangeable", scale.fix = TRUE, scale.value = 1)

## Beginning Cgee S-function, @(#) geeformula.q 4.13 98/01/27
## running glm to get initial regression estimate
## (Intercept)          age
## 0.64576265 -0.01201525

summary(resp_gee2)

##
## GEE: GENERALIZED LINEAR MODELS FOR DEPENDENT DATA
## gee S-function, version 4.13 modified 98/01/27 (1998)
##
## Model:
## Link:                               Logit
## Variance to Mean Relation: Binomial
## Correlation Structure:      Exchangeable
##
## Call:
## gee(formula = status ~ age, id = subject, data = RESPL, family = binomial,
##      corstr = "exchangeable", scale.fix = TRUE, scale.value = 1)
##
## Summary of Residuals:
##      Min       1Q   Median       3Q      Max
## -0.6256488 -0.5508736  0.3828310  0.4320996  0.5427142
##
##
## Coefficients:
##              Estimate Naive S.E.   Naive z Robust S.E.  Robust z
## (Intercept)  0.64576265 0.41037540  1.573590  0.39567302  1.632061
## age         -0.01201525 0.01135004 -1.058609  0.01173464 -1.023912
##
## Estimated Scale Parameter:  1
## Number of Iterations:  1
##
## Working Correlation
##      [,1]      [,2]      [,3]      [,4]
## [1,] 1.0000000 0.5279328 0.5279328 0.5279328
## [2,] 0.5279328 1.0000000 0.5279328 0.5279328
## [3,] 0.5279328 0.5279328 1.0000000 0.5279328
## [4,] 0.5279328 0.5279328 0.5279328 1.0000000

```

Table 10.7

```
formula3 <- status ~ age + month + baseline + centre + treatment + gender
resp_gee3 <- gee(formula3, data = RESPL, family = binomial, id = subject,
                 corstr = "exchangeable", scale.fix = TRUE, scale.value = 1)
```

```
## Beginning Cgee S-function, @(#) geeformula.q 4.13 98/01/27
```

```
## running glm to get initial regression estimate
```

```
##      (Intercept)          age          month
##      -0.74049186      -0.01818538      -0.06425079
##      baseline          centre2 treatmenttreatment
##      1.88406115      0.67230658      1.30055737
##      genderfemale
##      0.11936743
```

```
summary(resp_gee3)
```

```
##
```

```
## GEE: GENERALIZED LINEAR MODELS FOR DEPENDENT DATA
```

```
## gee S-function, version 4.13 modified 98/01/27 (1998)
```

```
##
```

```
## Model:
```

```
## Link:                      Logit
```

```
## Variance to Mean Relation: Binomial
```

```
## Correlation Structure:      Exchangeable
```

```
##
```

```
## Call:
```

```
## gee(formula = formula3, id = subject, data = RESPL, family = binomial,
```

```
##      corstr = "exchangeable", scale.fix = TRUE, scale.value = 1)
```

```
##
```

```
## Summary of Residuals:
```

```
##      Min      1Q      Median      3Q      Max
## -0.93336064 -0.29731249  0.08709546  0.32904495  0.85591042
```

```
##
```

```
##
```

```
## Coefficients:
```

```
##      Estimate Naive S.E.      Naive z Robust S.E.
## (Intercept) -0.73323775 0.51777213 -1.4161399 0.51901959
## age         -0.01841147 0.01256999 -1.4647157 0.01301183
## month       -0.06419992 0.08100426 -0.7925499 0.08150945
## baseline    1.87781449 0.34189674  5.4923439 0.35009816
## centre2     0.68091583 0.33960534  2.0050210 0.35676012
## treatmenttreatment 1.29222142 0.33562959  3.8501415 0.35053559
## genderfemale 0.13076523 0.41791342  0.3129003 0.44409802
```

```
##      Robust z
```

```
## (Intercept) -1.4127362
```

```
## age         -1.4149795
```

```
## month       -0.7876377
```

```
## baseline    5.3636799
```

```
## centre2     1.9086097
```

```
## treatmenttreatment 3.6864200
```

```
## genderfemale 0.2944513
```

```
##
```

```
## Estimated Scale Parameter: 1
## Number of Iterations: 2
##
## Working Correlation
##      [,1]      [,2]      [,3]      [,4]
## [1,] 1.000000 0.336918 0.336918 0.336918
## [2,] 0.336918 1.000000 0.336918 0.336918
## [3,] 0.336918 0.336918 1.000000 0.336918
## [4,] 0.336918 0.336918 0.336918 1.000000
```

Table 10.10

```

# install.packages("lme4")
library("lme4")

## Loading required package: Matrix
##
## Attaching package: 'Matrix'
## The following object is masked from 'package:tidyr':
##
##     expand

formula4 <- status ~ age + month + baseline + centre + treatment + gender + (1 | subject)
resp_glmer <- glmer(formula4, data = RESPL, family = binomial)
summary(resp_glmer)

## Generalized linear mixed model fit by maximum likelihood (Laplace
##   Approximation) [glmerMod]
##   Family: binomial   ( logit )
##   Formula: status ~ age + month + baseline + centre + treatment + gender +
##     (1 | subject)
##   Data: RESPL
##
##           AIC          BIC    logLik deviance df.resid
##        444.3        477.1    -214.2   428.3      436
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -2.8574 -0.3590  0.1427  0.3693  2.2393
##
## Random effects:
##   Groups Name            Variance Std.Dev.
##   subject (Intercept) 3.89      1.972
## Number of obs: 444, groups:  subject, 111
##
## Fixed effects:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -1.39580    0.83828  -1.665 0.095895 .
## age           -0.02546    0.02014  -1.264 0.206131
## month         -0.10133    0.12518  -0.810 0.418221
## baseline       3.07833    0.60271   5.107 3.26e-07 ***
## centre2        1.04669    0.54784   1.911 0.056059 .
## treatmenttreatment 2.16324    0.55643   3.888 0.000101 ***
## genderfemale    0.20250    0.67270   0.301 0.763394
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) age      month  baseln centr2 trtmnt
## age          -0.670
## month        -0.356  0.009
## baseline     -0.349  0.015 -0.041
## centre2      -0.078 -0.223 -0.015 -0.150

```

```
## trtmnttrtmn -0.395 -0.050 -0.031 0.301 0.058
## genderfemal -0.008 -0.263 -0.003 0.102 -0.147 0.219
```

Table 10.2: Data from a Clinical Trial of Patients Suffering from Epilepsy

```
library("HSAUR3")
EPIL <- epilepsy # (data are readily given in long form)
glimpse(EPIL); head(EPIL)
```

```
## Observations: 236
## Variables: 6
## $ treatment    <fct> placebo, placebo, placebo, placebo, placebo, plac...
## $ base         <int> 11, 11, 11, 11, 11, 11, 11, 11, 6, 6, 6, 6, 8, 8,...
## $ age          <int> 31, 31, 31, 31, 30, 30, 30, 30, 25, 25, 25, 25, 3...
## $ seizure.rate <int> 5, 3, 3, 3, 3, 5, 3, 3, 2, 4, 0, 5, 4, 4, 1, 4, 7...
## $ period       <ord> 1, 2, 3, 4, 1, 2, 3, 4, 1, 2, 3, 4, 1, 2, 3, 4, 1...
## $ subject      <fct> 1, 1, 1, 1, 2, 2, 2, 2, 3, 3, 3, 3, 4, 4, 4, 4, 5...
```

	treatment	base	age	seizure.rate	period	subject
## 1	placebo	11	31	5	1	1
## 110	placebo	11	31	3	2	1
## 112	placebo	11	31	3	3	1
## 114	placebo	11	31	3	4	1
## 2	placebo	11	30	3	1	2
## 210	placebo	11	30	5	2	2

```
# convert data to WIDE form:
# (sep: set valid names month0:month4 here! otherwise just 0:4)
```

```
EPI <- EPIL %>%
  spread(key = period, value = seizure.rate, sep = "") %>%
  arrange(subject)
```

```
glimpse(EPI); head(EPI)
```

```
## Observations: 59
## Variables: 8
## $ treatment    <fct> placebo, placebo, placebo, placebo, placebo, placebo...
## $ base         <int> 11, 11, 6, 8, 66, 27, 12, 52, 23, 10, 52, 33, 18, 42...
## $ age          <int> 31, 30, 25, 36, 22, 29, 31, 42, 37, 28, 36, 24, 23, ...
## $ subject      <fct> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 1...
## $ period1      <int> 5, 3, 2, 4, 7, 5, 6, 40, 5, 14, 26, 12, 4, 7, 16, 11...
## $ period2      <int> 3, 5, 4, 4, 18, 2, 4, 20, 6, 13, 12, 6, 4, 9, 24, 0,...
## $ period3      <int> 3, 3, 0, 1, 9, 8, 0, 23, 6, 6, 6, 8, 6, 12, 10, 0, 3...
## $ period4      <int> 3, 3, 5, 4, 21, 7, 2, 12, 5, 0, 22, 4, 2, 14, 9, 5, ...
```

	treatment	base	age	subject	period1	period2	period3	period4
## 1	placebo	11	31	1	5	3	3	3
## 2	placebo	11	30	2	3	5	3	3
## 3	placebo	6	25	3	2	4	0	5
## 4	placebo	8	36	4	4	4	1	4
## 5	placebo	66	22	5	7	18	9	21
## 6	placebo	27	29	6	5	2	8	7

```
# correct the baseline when converting to long form:
```

```
EPIL <- gather(EPI, key = periods, value = seizure.rate, period1:period4) %>%
```



```
mutate( base2 = base/4, week = as.integer(substr(weeks, 7, 7)) ) %>%
  arrange(subject)
```

```
glimpse(EPIL); head(EPIL)
```

```
## Observations: 236
## Variables: 8
## $ treatment    <fct> placebo, placebo, placebo, placebo, placebo, plac...
## $ base         <int> 11, 11, 11, 11, 11, 11, 11, 11, 6, 6, 6, 6, 8, 8,...
## $ age          <int> 31, 31, 31, 31, 30, 30, 30, 30, 25, 25, 25, 25, 3,...
## $ subject      <fct> 1, 1, 1, 1, 2, 2, 2, 2, 3, 3, 3, 3, 4, 4, 4, 4, 5...
## $ periods      <chr> "period1", "period2", "period3", "period4", "peri...
## $ seizure.rate <int> 5, 3, 3, 3, 3, 5, 3, 3, 2, 4, 0, 5, 4, 4, 1, 4, 7...
## $ base2        <dbl> 2.75, 2.75, 2.75, 2.75, 2.75, 2.75, 2.75, 2.75, 2.75, 1...
## $ week         <int> 1, 2, 3, 4, 1, 2, 3, 4, 1, 2, 3, 4, 1, 2, 3, 4, 1...

##   treatment base age subject periods seizure.rate base2 week
## 1 placebo   11  31      1 period1          5  2.75    1
## 2 placebo   11  31      1 period2          3  2.75    2
## 3 placebo   11  31      1 period3          3  2.75    3
## 4 placebo   11  31      1 period4          3  2.75    4
## 5 placebo   11  30      2 period1          3  2.75    1
## 6 placebo   11  30      2 period2          5  2.75    2
```

Figure 10.3

```
# Convert data to long form, including the baseline;
# correct the baseline first:
EPIO <- EPI %>%
  mutate(base2 = base/4)

EPILO <- gather(EPIO, key = periods, value = seizure.rate, base2, period1:period4) %>%
  arrange(subject)

glimpse(EPILO); head(EPILO)

## Observations: 295
## Variables: 6
## $ treatment    <fct> placebo, placebo, placebo, placebo, placebo, plac...
## $ base         <int> 11, 11, 11, 11, 11, 11, 11, 11, 11, 11, 6, 6, 6, ...
## $ age          <int> 31, 31, 31, 31, 31, 30, 30, 30, 30, 30, 25, 25, 2...
## $ subject      <fct> 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 3, 3, 3, 3, 3, 4, 4...
## $ periods      <chr> "base2", "period1", "period2", "period3", "period...
## $ seizure.rate <dbl> 2.75, 5.00, 3.00, 3.00, 3.00, 2.75, 3.00, 5.00, 3...

##   treatment base age subject periods seizure.rate
## 1 placebo   11  31      1   base2         2.75
## 2 placebo   11  31      1 period1         5.00
## 3 placebo   11  31      1 period2         3.00
## 4 placebo   11  31      1 period3         3.00
## 5 placebo   11  31      1 period4         3.00
## 6 placebo   11  30      2   base2         2.75

# Mark some of the outliers:
# https://stackoverflow.com/questions/33524669/labeling-outliers-of-boxplots-in-r
is_outlier <- function(x) {
  # changed the upper coeff from 1.5 to 3.5 (only more extreme observations highlighted)
  return(x < quantile(x, 0.25) - 1.5 * IQR(x) | x > quantile(x, 0.75) + 3.5 * IQR(x))
}

# use subject id (not seizure.rate itself!) for identification:
EPILO <- EPILO %>%
  mutate(outlier = ifelse(is_outlier(seizure.rate), subject, as.numeric(NA)))

p1 <- ggplot(EPILO, aes(x = factor(periods), y = seizure.rate, fill = treatment))
p2 <- p1 + geom_boxplot() + geom_text(aes(label = outlier), na.rm = TRUE, hjust = -0.6)
p3 <- p2 + facet_grid(treatment ~ ., labeller = label_parsed)
p4 <- p3 + theme_bw() + theme(panel.grid.major = element_blank(),
                             panel.grid.minor = element_blank())
p5 <- p4 + theme(legend.position = "none")
p6 <- p5 + scale_x_discrete(name = "",
                           labels = c("Baseline", "Time 1", "Time 2", "Time 3", "Time 4"))
p7 <- p6 + scale_y_continuous(name = "Seizure rate (count)")
p8 <- p7 + scale_fill_grey(start = 1, end = 1) # B&W for the book
p7; p8
```

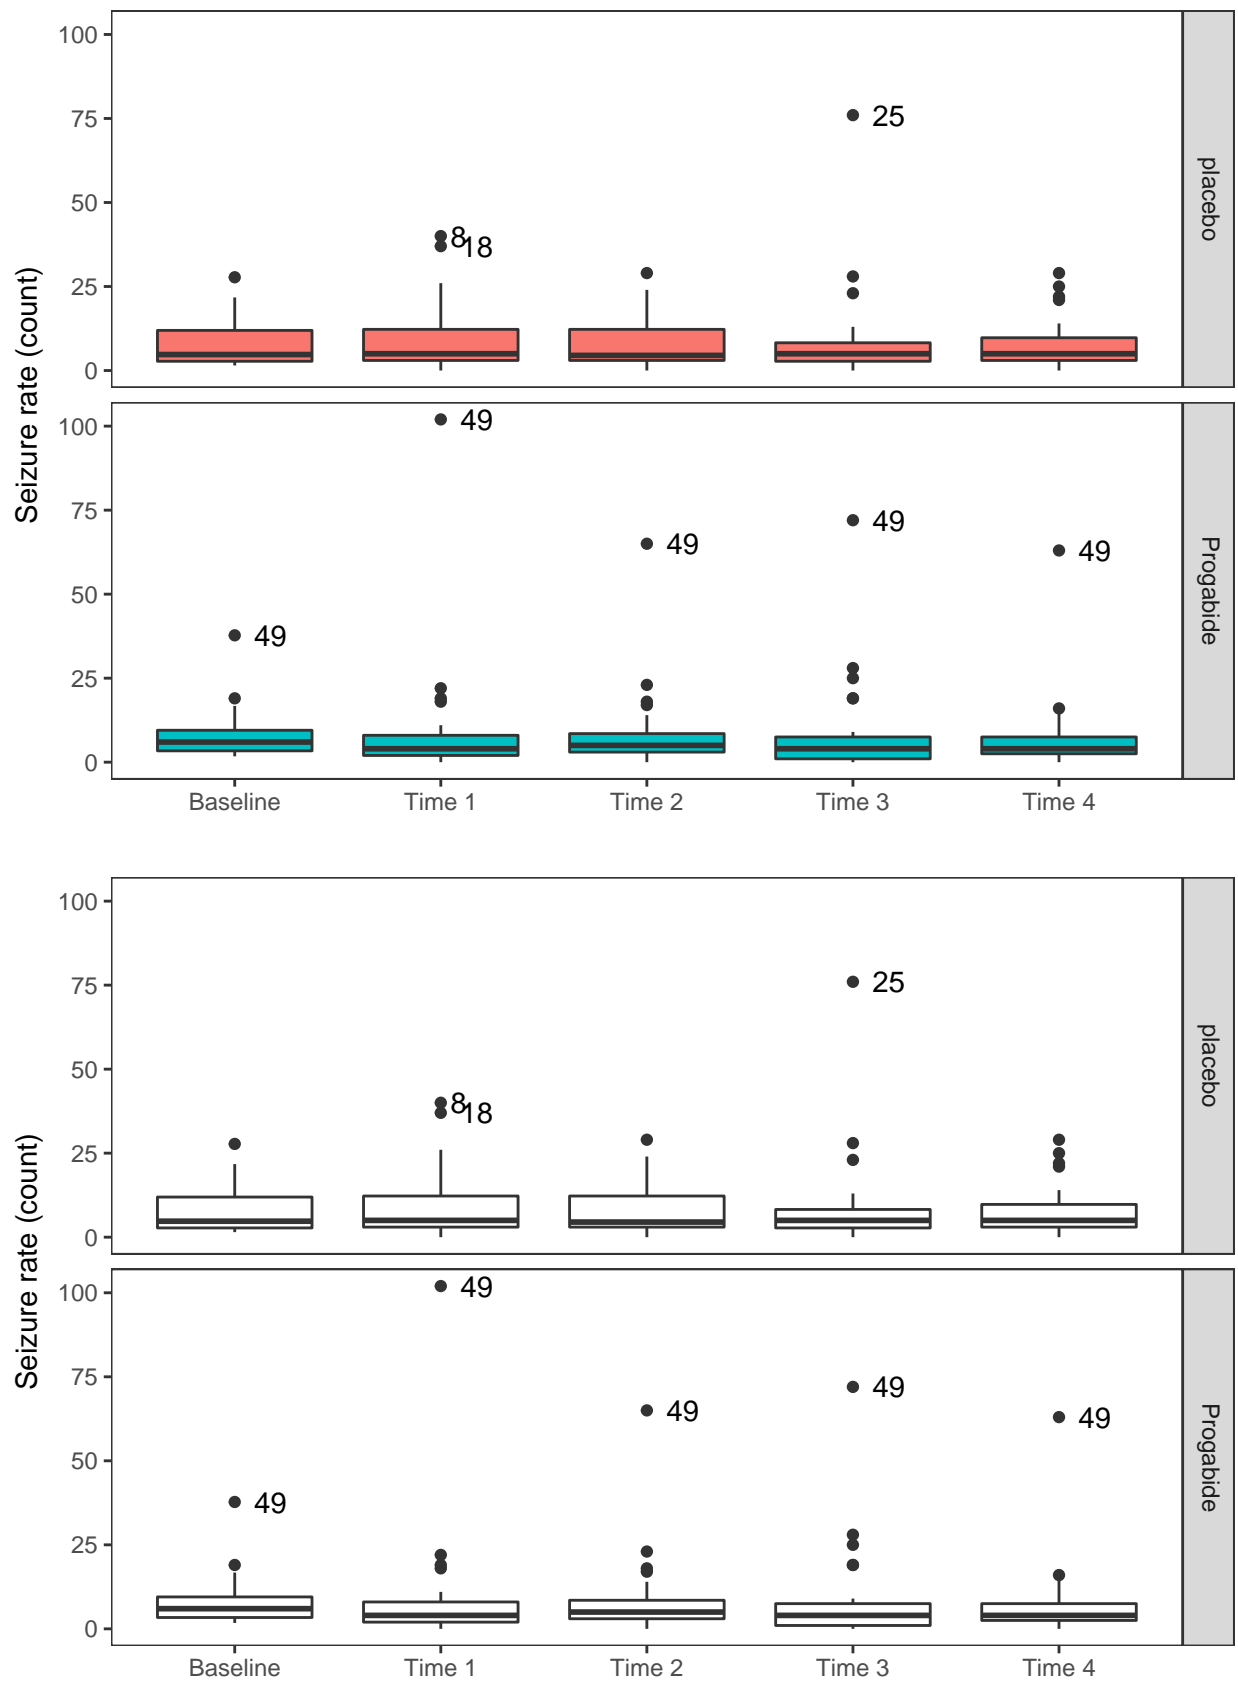


Table 10.8

```
t1 <- tapply(EPILO$seizure.rate, EPILO$period, mean)
t2 <- tapply(EPILO$seizure.rate, EPILO$period, var)
t1 <- as.data.frame(t1)
t2 <- as.data.frame(t2)

tab108 <- cbind(t1, t2)
names(tab108) = c("mean", "variance")
tab108
```

```
##           mean  variance
## base2    7.805085  45.14885
## period1  8.949153 220.08358
## period2  8.355932 103.78492
## period3  8.440678 200.18177
## period4  7.305085  93.11222
```

Table 10.9

```

formula5 <- seizure.rate ~ week + age + log(base2) + treatment
epil_gee1 <- gee(formula5, data = EPIL, family = "poisson", id = subject,
                 corstr = "exchangeable", scale.fix = FALSE)

## Beginning Cgee S-function, @(#) geeformula.q 4.13 98/01/27
## running glm to get initial regression estimate

##      (Intercept)           week           age
##      -0.84405973      -0.05872397      0.01932016
##      log(base2) treatmentProgabide
##      1.22082626      -0.03005351

summary(epil_gee1)

##
## GEE: GENERALIZED LINEAR MODELS FOR DEPENDENT DATA
## gee S-function, version 4.13 modified 98/01/27 (1998)
##
## Model:
## Link:                               Logarithm
## Variance to Mean Relation: Poisson
## Correlation Structure:      Exchangeable
##
## Call:
## gee(formula = formula5, id = subject, data = EPIL, family = "poisson",
##      corstr = "exchangeable", scale.fix = FALSE)
##
## Summary of Residuals:
##      Min      1Q      Median      3Q      Max
## -15.5346018  -3.2421580  -0.3594229   1.4768610  60.2253586
##
##
## Coefficients:
##              Estimate Naive S.E.   Naive z Robust S.E.
## (Intercept)   -0.87530135  0.52946461 -1.6531820  0.473365597
## week          -0.05872335  0.03459127 -1.6976352  0.034997602
## age            0.01999638  0.01250451  1.5991331  0.009792718
## log(base2)     1.22469907  0.10554688 11.6033662  0.155673659
## treatmentProgabide -0.02421150  0.15532727 -0.1558741  0.191087900
##              Robust z
## (Intercept)   -1.8491022
## week          -1.6779251
## age            2.0419638
## log(base2)     7.8670925
## treatmentProgabide -0.1267035
##
## Estimated Scale Parameter:  4.822177
## Number of Iterations:  2
##
## Working Correlation
##      [,1]      [,2]      [,3]      [,4]
## [1,] 1.0000000  0.3986404  0.3986404  0.3986404

```

```
## [2,] 0.3986404 1.0000000 0.3986404 0.3986404  
## [3,] 0.3986404 0.3986404 1.0000000 0.3986404  
## [4,] 0.3986404 0.3986404 0.3986404 1.0000000
```

Table 10.11

```

formula6 <- seizure.rate ~ week + age + log(base2) + treatment + (1 | subject)
epil_glmml <- glmer(formula6, data = EPIL, family = "poisson")

## Warning in checkConv(attr(opt, "derivs"), opt$par, ctrl = control
## $checkConv, : Model failed to converge with max|grad| = 0.00241629 (tol =
## 0.001, component 1)

summary(epil_glmml)

## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: poisson ( log )
## Formula:
## seizure.rate ~ week + age + log(base2) + treatment + (1 | subject)
## Data: EPIL
##
##           AIC          BIC    logLik deviance df.resid
##      1346.2      1367.0    -667.1   1334.2       230
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.2981 -0.8756 -0.0889  0.6043  7.2421
##
## Random effects:
## Groups Name Variance Std.Dev.
## subject (Intercept) 0.268 0.5177
## Number of obs: 236, groups: subject, 59
##
## Fixed effects:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -0.17533    0.43865  -0.400  0.68937
## week          -0.05872    0.02018  -2.910  0.00362 **
## age            0.01073    0.01217   0.882  0.37800
## log(base2)     1.02527    0.10115  10.136 < 2e-16 ***
## treatmentProgabide -0.31998    0.15060  -2.125  0.03362 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) week    age    lg(b2)
## week          -0.112
## age          -0.872  0.000
## log(base2)   -0.545  0.000  0.161
## trtmntPrgrbd -0.243  0.000  0.110 -0.043
## convergence code: 0
## Model failed to converge with max|grad| = 0.00241629 (tol = 0.001, component 1)

```

Table 10.12

```
formula7 <- seizure.rate ~ week + age + log(base2) + treatment + (1 + week | subject)
epil_glmm2 <- glmer(formula7, data = EPIL, family = "poisson")
```

```
## Warning in checkConv(attr(opt, "derivs"), opt$par, ctrl = control
## $checkConv, : Model failed to converge with max|grad| = 0.0012166 (tol =
## 0.001, component 1)
```

```
summary(epil_glmm2)
```

```
## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: poisson ( log )
## Formula: seizure.rate ~ week + age + log(base2) + treatment + (1 + week |
## subject)
## Data: EPIL
##
##           AIC          BIC    logLik deviance df.resid
##    1330.4    1358.1    -657.2   1314.4      228
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -2.9840 -0.7807 -0.1003  0.5385  6.3781
##
## Random effects:
##   Groups Name            Variance Std.Dev. Corr
##   subject (Intercept) 0.4134    0.6429
##   week              0.0213    0.1459  -0.60
## Number of obs: 236, groups: subject, 59
##
## Fixed effects:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -0.18761    0.44351  -0.423   0.6723
## week          -0.05135    0.03264  -1.573   0.1157
## age           0.01001    0.01233   0.812   0.4166
## log(base2)     1.02520    0.10085  10.165 <2e-16 ***
## treatmentProgabide -0.31674    0.15018  -2.109   0.0349 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) week    age    lg(b2)
## week          -0.136
## age          -0.864 -0.051
## log(base2)   -0.537 -0.019  0.162
## trtmntPrghd -0.243 -0.004  0.113 -0.042
## convergence code: 0
## Model failed to converge with max|grad| = 0.0012166 (tol = 0.001, component 1)
```